

```
In [1]: !pip install pandas numpy matplotlib seaborn scikit-learn

Requirement already satisfied: pandas in /home/sargam/.conda/envs/notebook
s/lib/python3.10/site-packages (2.3.3)
Requirement already satisfied: numpy in /home/sargam/.conda/envs/notebook
s/lib/python3.10/site-packages (2.2.6)
Requirement already satisfied: matplotlib in /home/sargam/.conda/envs/note
books/lib/python3.10/site-packages (3.10.6)
Requirement already satisfied: seaborn in /home/sargam/.conda/envs/noteboo
ks/lib/python3.10/site-packages (0.13.2)
Requirement already satisfied: scikit-learn in /home/sargam/.conda/envs/no
tebooks/lib/python3.10/site-packages (1.7.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /home/sargam/.con
da/envs/notebooks/lib/python3.10/site-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /home/sargam/.conda/envs/no
tebooks/lib/python3.10/site-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /home/sargam/.conda/envs/
notebooks/lib/python3.10/site-packages (from pandas) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in /home/sargam/.conda/en
vs/notebooks/lib/python3.10/site-packages (from matplotlib) (1.3.2)
Requirement already satisfied: cycler>=0.10 in /home/sargam/.conda/envs/no
tebooks/lib/python3.10/site-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /home/sargam/.conda/en
vs/notebooks/lib/python3.10/site-packages (from matplotlib) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /home/sargam/.conda/en
vs/notebooks/lib/python3.10/site-packages (from matplotlib) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /home/sargam/.conda/en
vs/notebooks/lib/python3.10/site-packages (from matplotlib) (25.0)
Requirement already satisfied: pillow>=8 in /home/sargam/.conda/envs/note
books/lib/python3.10/site-packages (from matplotlib) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /home/sargam/.conda/en
vs/notebooks/lib/python3.10/site-packages (from matplotlib) (3.2.5)
Requirement already satisfied: scipy>=1.8.0 in /home/sargam/.conda/envs/no
tebooks/lib/python3.10/site-packages (from scikit-learn) (1.15.3)
Requirement already satisfied: joblib>=1.2.0 in /home/sargam/.conda/envs/n
otebooks/lib/python3.10/site-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /home/sargam/.cond
a/envs/notebooks/lib/python3.10/site-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: six>=1.5 in /home/sargam/.conda/envs/note
books/lib/python3.10/site-packages (from python-dateutil>=2.8.2->pandas) (1.
17.0)
```

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [5]: data = pd.read_csv('/home/sargam/things/college/daa/lp3/ml/uber/uber.csv')
```

```
In [6]: print("Dataset shape:", data.shape)
print(data.info())
print(data.head())
```

```
Dataset shape: (200000, 9)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        200000 non-null   int64  
 1   key               200000 non-null   object  
 2   fare_amount       200000 non-null   float64 
 3   pickup_datetime   200000 non-null   object  
 4   pickup_longitude  200000 non-null   float64 
 5   pickup_latitude   200000 non-null   float64 
 6   dropoff_longitude 199999 non-null   float64 
 7   dropoff_latitude  199999 non-null   float64 
 8   passenger_count   200000 non-null   int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
None
      Unnamed: 0          key  fare_amount \
0    24238194  2015-05-07 19:52:06.0000003    7.5
1    27835199  2009-07-17 20:04:56.0000002    7.7
2    44984355  2009-08-24 21:45:00.0000001   12.9
3    25894730  2009-06-26 08:22:21.0000001    5.3
4    17610152  2014-08-28 17:47:00.000000188   16.0

      pickup_datetime  pickup_longitude  pickup_latitude \
0  2015-05-07 19:52:06 UTC           -73.999817     40.738354
1  2009-07-17 20:04:56 UTC           -73.994355     40.728225
2  2009-08-24 21:45:00 UTC           -74.005043     40.740770
3  2009-06-26 08:22:21 UTC           -73.976124     40.790844
4  2014-08-28 17:47:00 UTC           -73.925023     40.744085

      dropoff_longitude  dropoff_latitude  passenger_count
0           -73.999512      40.723217                  1
1           -73.994710      40.750325                  1
2           -73.962565      40.772647                  1
3           -73.965316      40.803349                  3
4           -73.973082      40.761247                  5
```

```
In [7]: data.isnull().sum()
```


```
Out[7]: Unnamed: 0      0
key          0
fare_amount  0
pickup_datetime 0
pickup_longitude 0
pickup_latitude 0
dropoff_longitude 1
dropoff_latitude 1
passenger_count 0
dtype: int64
```

```
In [8]: data['pickup_datetime'] = pd.to_datetime(data['pickup_datetime'], utc=True)
```

```
In [9]: data.dropna(inplace=True)
```

```
In [10]: plt.figure(figsize=(8, 4))
sns.boxplot(data['fare_amount'])
```

```
plt.title('Boxplot for Fare Amount')
plt.show()
```



```
In [11]: import matplotlib.pyplot as plt
import seaborn as sns


# Plot the distribution and boxplot of fare_amount to identify outliers
plt.figure(figsize=(14, 6))
```

```
Out[11]: <Figure size 1400x600 with 0 Axes>
<Figure size 1400x600 with 0 Axes>
```

```
In [12]: plt.subplot(1, 2, 1)
sns.histplot(data['fare_amount'], bins=50, kde=True)
plt.title('Fare Amount Distribution')

plt.subplot(1, 2, 2)
sns.boxplot(x=data['fare_amount'])
plt.title('Fare Amount Boxplot')

plt.show()
```



```
In [13]: plt.figure(figsize=(14, 8))


plt.subplot(2, 2, 1)
sns.boxplot(x=data['pickup_longitude'])
plt.title('Pickup Longitude Boxplot')

plt.subplot(2, 2, 2)
sns.histplot(data['pickup_longitude'], bins=50, kde=True)
plt.title('Pickup Longitude Distribution')

plt.subplot(2, 2, 3)
sns.boxplot(x=data['pickup_latitude'])
plt.title('Pickup Latitude Boxplot')


plt.subplot(2, 2, 4)
sns.histplot(data['pickup_latitude'], bins=50, kde=True)
plt.title('Pickup Latitude Distribution')

plt.tight_layout()
plt.show()
```



```
In [14]: # Remove outliers outside 1st and 99th percentile
low_fare = data['fare_amount'].quantile(0.01)
high_fare = data['fare_amount'].quantile(0.99)
data = data[(data['fare_amount'] >= low_fare) & (data['fare_amount'] <= h]
```

```
In [15]: # Step 5: Correlation analysis
numeric_features = data.select_dtypes(include=[np.number])
corr = numeric_features.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



```
In [16]: # Step 6: Feature Engineering
# Convert datetime to numerical features such as hour, day, weekday
data['pickup_hour'] = data['pickup_datetime'].dt.hour
data['pickup_day'] = data['pickup_datetime'].dt.day
data['pickup_weekday'] = data['pickup_datetime'].dt.weekday
```

```
In [17]: # Step 7: Select Features and Target
# Drop columns that are not useful or contain identifiers
X = data.drop(columns=['fare_amount', 'pickup_datetime', 'key', 'Unnamed: 0'])
y = data['fare_amount']
```

```
In [18]: # Step 8: Split dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
```

```
In [19]: # Step 9: Implement Linear Regression Model
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
y_pred_lr = lr_model.predict(X_test)
```

```
In [20]: # Step 10: Implement Random Forest Regression Model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
y_pred_rf = rf_model.predict(X_test)
```


```
In [21]: # Step 11: Evaluate models with RMSE and R2 Score
def evaluate_model(true, pred, model_name):
    rmse = np.sqrt(mean_squared_error(true, pred))
```

```
r2 = r2_score(true, pred)
print(f"{model_name} RMSE: {rmse:.4f}")
print(f"{model_name} R2 Score: {r2:.4f}")
```

```
In [22]: evaluate_model(y_test, y_pred_lr, "Linear Regression")
evaluate_model(y_test, y_pred_rf, "Random Forest Regression")
```

```
Linear Regression RMSE: 8.0621
Linear Regression R2 Score: 0.0004
Random Forest Regression RMSE: 3.5102
Random Forest Regression R2 Score: 0.8105
```

```
In [23]: # Optional: Visualize actual vs predicted for the best model
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred_rf, alpha=0.3)
plt.xlabel("Actual Fare Amount")
plt.ylabel("Predicted Fare Amount")
plt.title("Random Forest Regression: Actual vs Predicted")
plt.show()
```



```
In [ ]:
```